

# Final Project: Walmart Data Exploration

*Stephen Kappel (spk2131), Mayank Misra (mm3557), Mandeep Singh (ms4826)*

*Due: December 5, 2015*

## Introduction

For our project, we chose to use sales data provided by Walmart as part of the *Walmart Recruiting - Store Sales Forecasting* competition on Kaggle in 2014. (See <https://www.kaggle.com/c/walmart-recruiting-store-sales-forecasting>.) The dataset contains weekly sales for 45 Walmart stores from 2/5/2010 to 11/1/2012. The sales are broken out by store and by department. Other attributes provide further context to the sales numbers:

- *Economic indicators*: CPI, unemployment rate, fuel prices
- *Store-specific attributes*: store size, store type, markdowns (indicator variables for five types of markdowns)
- *Other*: temperature, holiday (indicator variable)

In our analysis, sales is the outcome variable of primary interest. We aim to understand what factors impact sales. In particular, we try to answer:

1. How do economic factors – over which Walmart has no control – affect Walmart’s sales?
2. How well can we predict sales? In predicting sales, which predictors are most significant?

## Data Preparation

Kaggle provides the data in the form of three CSVs. We a data.frames that combines the three source CSVs together. `store.level` contains one row per store per week. (The original data had an additional “department” level of detail that we have removed.)

```
require(plyr)

get.stores <- function(){
  stores <- read.csv('../data/stores.csv')
  stores$Store <- as.factor(stores$Store)
  return(stores)
}

get.features <- function(){
  features <- read.csv('../data/features.csv')
  features$Store <- as.factor(features$Store)
  features$Date <- as.Date(features$Date)
  return(features)
}

get.train <- function(){
  train <- read.csv('../data/train.csv')
  train$Store <- as.factor(train$Store)
  train$Dept <- as.factor(train$Dept)
  train$Date <- as.Date(train$Date)
```

```

    return(train)
}

get.store.level <- function(){
  # construct a data.frame with detail only down to the store level (no departments)
  train <- get.train()
  train <- ddply(train, c('Store', 'Date'), summarize, Sales=sum(Weekly_Sales))
  stores <- get.stores()
  features <- get.features()
  merged <- merge(train, stores, by='Store')
  merged <- merge(merged, features, by=c('Store', 'Date'))
  merged$Sales_Millions <- merged$Sales / 10e6
  merged <- rename(merged, c('Type'='Store_Type', 'Size'='Store_Size'))
  return(merged)
}

store.level <- get.store.level()

```

As we do our analysis, we often find it helpful to normalize a column within the context of a store. This allows for more meaningful comparisons. Because this is a common operation, we define a function to add a normalized column to a given data.frame for a specified column/attribute. To start, we add a normalized sales column to the `store.level` data.frame.

```

add.normalized.col <- function(df, col.name){
  # df: data.frame (usually store.level) to add normalized column to
  # col.name: the name of the column in df which should be normalized
  # the normalized column is named as 'Norm_[col.name]'
  df$temp <- df[, col.name]
  store.ply <- ddply(df, c('Store'), summarize, Mean_X=mean(temp),
                    SD_X=sd(temp))
  df <- merge(df, store.ply, by=c('Store'))
  df$Norm_X <- (df[,col.name] - df$Mean_X) / df$SD_X
  df$Mean_X <- NULL
  df$SD_X <- NULL
  df$temp <- NULL
  df <- rename(df, c('Norm_X'=paste('Norm_', col.name, sep='')))
  return(df)
}

store.level <- add.normalized.col(store.level, 'Sales')

```

We also commonly want to get sales aggregates after grouping and/or filtering by some attribute, so we define a function for this operation.

```

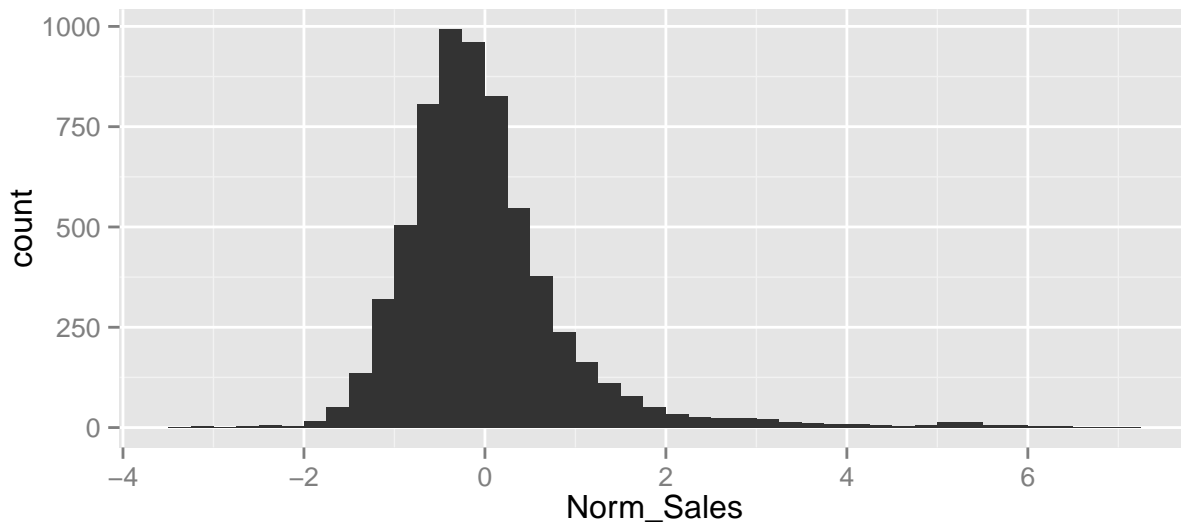
get.aggr.sales <- function(df, aggr.dims, filter=NULL){
  # df: data.frame (usually store.level)
  # aggr.dims: a list of column names to aggregate on
  # filter: NULL or c(column name to filter, value to filter to)
  if(!is.null(filter)){
    sub.df <- subset(df, df[,filter[1]] == filter[2])
  } else{
    sub.df <- df
  }
}

```

```
grouped <- ddply(sub.df, aggr.dims, summarize, Sales=sum(Sales),
                 Sales_Millions=sum(Sales_Millions))
return(grouped)
}
```

Let's take a look at the normalized distribution of store-level weekly sales. We hope this is roughly normal as it will allow us to apply some statistical tests with more confidence in the coming analysis.

```
require(ggplot2)
ggplot(data=store.level, aes(x=Norm_Sales)) + geom_histogram(binwidth=0.25)
```




The distribution appears roughly symmetric and normal, although there is a light long tail on the positive side.

## Impact of economic factors

In this section we explore how economic factors (unemployment, CPI, and fuel price) relate to Walmart's sales. It's not immediately obvious how we should expect Walmart's sales to vary with economic conditions. While poor economic conditions hurt most retailers, Walmart is known for low prices. Could bad economic times drive more people to shop at Walmart?


We start by creating a boxplot of sales by unemployment rate.

```
ggplot(data=store.level, aes(x=Unemployment, y=Sales_Millions,
                             group=round_any(Unemployment, 1.0))) + geom_boxplot()
```




This plot don't show much of a relationship, but we haven't done any normalization, so this probably isn't a meaningful representation of the relationship. We normalize sales and unemployment rate by store and create an update boxplot.

```
store.level <- add.normalized.col(store.level, 'Unemployment')
ggplot(data=store.level, aes(x=Norm_Unemployment, y=Norm_Sales,
  group=round_any(Norm_Unemployment, 0.25))) + geom_boxplot()
```




The same data in a scatter plot with LOESS (blue) and GAM (red) curves overlaid:

```
ggplot(data=store.level, aes(x=Norm_Unemployment, y=Norm_Sales)) +
  geom_point(alpha = 1/5) + geom_smooth(method='loess', color='blue') +
  geom_smooth(method='gam', color='red')
```



Within the -1 to 1 normalized unemployment range, a pattern seems to have emerged. There appears to be below-average sales when there is above-average unemployment. Let's create a pair of histograms showing number of weeks by normalized sales - one histogram for weeks with below average unemployment and one histogram for weeks with above average unemployment. We want to see if these distributions are noticeably different. We also want to see the shape of these distributions to see if the data is normally distributed and a t-test would be reasonable.

```
store.level$Unemployment_Category <- ifelse(
  store.level$Norm_Unemployment < 0, 'Below Average', 'Above Average')
ggplot(data=store.level, aes(x=Norm_Sales)) + geom_histogram(binwidth=0.25) +
  facet_wrap(~Unemployment_Category)
```



The relationship between unemployment and sales isn't too clear from this plot, but the populations do appear to be roughly normal. We now run an F-test to check if the populations have equal variances (as assumed by a two-sample t-test).

```

above <- subset(store.level, Unemployment_Category == 'Above Average')
below <- subset(store.level, Unemployment_Category == 'Below Average')
var.test(above$Norm_Sales, below$Norm_Sales)

```

```

##
## F test to compare two variances
##
## data:  above$Norm_Sales and below$Norm_Sales
## F = 1.2265, num df = 3477, denom df = 2956, p-value = 9.068e-09
## alternative hypothesis: true ratio of variances is not equal to 1
## 95 percent confidence interval:
##  1.144162 1.314409
## sample estimates:
## ratio of variances
##           1.226456

```

We reject the null hypothesis that the populations have equal variance, so we proceed with Welch t-test for populations with unequal variances. The null hypothesis of this t-test is that the mean normalized sales during weeks with above average unemployment rates is equal to the mean normalized sales during weeks with below average unemployment rates.

```

t.test(above$Norm_Sales, below$Norm_Sales, var.equal=FALSE)

```

```

##
## Welch Two Sample t-test
##
## data:  above$Norm_Sales and below$Norm_Sales
## t = -6.831, df = 6409.719, p-value = 9.209e-12
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  -0.2166014 -0.1200037
## sample estimates:
##  mean of x   mean of y
## -0.07733810  0.09096447

```

We can confidently reject the null hypothesis. Now, we drop our assumption of normality, and perform a non-parametric test to see determine if sales for weeks with above-average unemployment are drawn from a different distribution than the sales for weeks with below-average unemployment. We apply the Mann-Whitney-Wilcoxon procedure below.

```

wilcox.test(Norm_Sales ~ Unemployment_Category, data=store.level)

```

```

##
## Wilcoxon rank sum test with continuity correction
##
## data:  Norm_Sales by Unemployment_Category
## W = 4330099, p-value < 2.2e-16
## alternative hypothesis: true location shift is not equal to 0

```

From all the results above, we can conclude that unemployment rate is not independent of Walmart sales.

We repeat the same procedures for CPI and fuel price. We hide the code and output to save space, but we summarize the p-values from the tests in the table below:

Economic Factor	F-test 95% CI	t-test 95% CI	Mann-Whitney-Wilcoxon p-value
Unemployment	(1.14, 1.31)	(-0.22, -0.12)	2.2e-16
CPI	(0.83, 0.95)	(0.11, 0.21)	2.2e-16
Fuel price	(0.34, 0.39)	(-0.15, -0.04)	3.9e-4

All three variables tell a similar story. When external economic factors are unfavorable (high unemployment, low CPI, and high fuel prices), Walmart sales are not as strong as they are during times when external economic factors are favorable.


## Sales predictions

With the goal of creating a model to predict store-level normalized sales for a given week and store, we build a set of features and then compare the performance of a few different models.


### Features

In order to make good predictions about sales, we need to consider much more than just the economic factors we've considered so far. For example, the next couple plots show that dates will also be important in our prediction; we observe seasonality and trending patterns in the timeseries.


```
# plot sales by week
sales.by.week <- get.aggr.sales(store.level, c('Date'))
ggplot(data=sales.by.week, aes(x=Date, y=Sales_Millions)) + geom_line()
```



```
# plot sales by week by store type
sales.by.type <- get.aggr.sales(store.level, c('Store_Type', 'Date'))
ggplot(data=sales.by.type, aes(x=Date, y=Sales_Millions, group=Store_Type,
                               color=Store_Type)) + geom_line()
```



```
# by store for type C
sales.by.store <- get.aggr.sales(store.level, c('Store', 'Date'), c('Store_Type', 'C'))
ggplot(data=sales.by.store, aes(x=Date, y=Sales_Millions, group=Store, color=Store)) +
  geom_line() + guides(color=FALSE)
```



To capture some of the seasonality and trending nature of the data, we add a month-of-year factor feature and we generate “lag” features that represent what the normalized sales were 1, 2, 52, and 53 weeks ago. We also create a feature that represents the difference in normalized sales between last week and 53 weeks ago.

```
require(lubridate)
store.level$Month <- as.factor(month(store.level$Date))

add.lag.column <- function(df, weeks){
  # df: data.frame (usually store.level) to which lag column will be added
  # weeks: an integer representing the number of weeks ago that a lag column
  #         will be created for
```



```

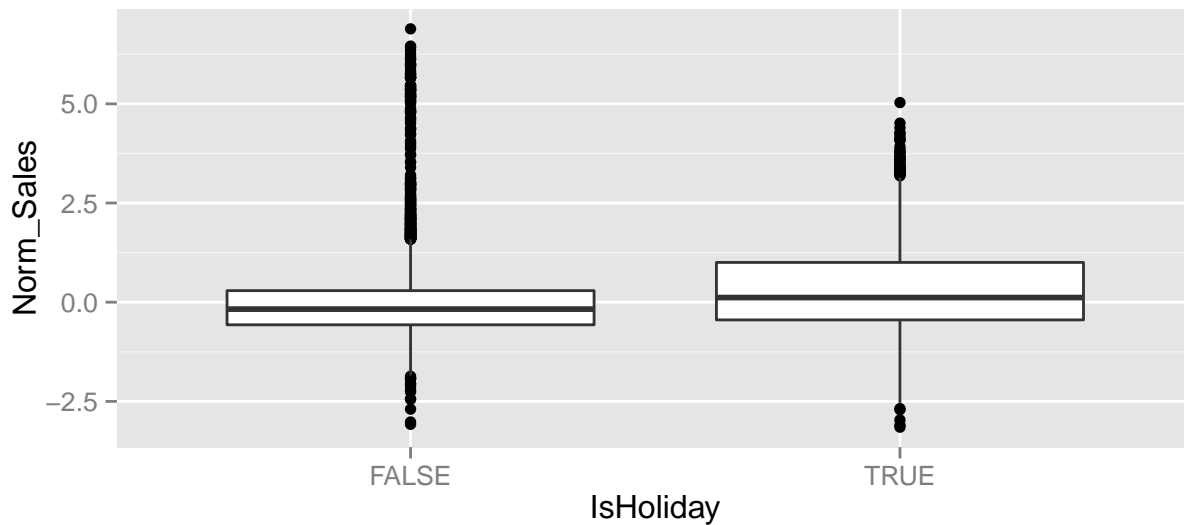
# The added lag column is named 'Lag_[weeks]_Sales'
shifted <- df[, c('Store', 'Date', 'Norm_Sales')]
shifted$Date <- shifted$Date + (weeks * 7)
shifted <- rename(shifted, c('Norm_Sales' = paste('Lag_', weeks, '_Sales', sep='')))
df <- merge(df, shifted, c('Store', 'Date'), all.x = TRUE, all.y = FALSE)
return(df)
}

store.level <- add.lag.column(store.level, 1)
store.level <- add.lag.column(store.level, 2)
store.level <- add.lag.column(store.level, 3)
store.level <- add.lag.column(store.level, 4)
store.level <- add.lag.column(store.level, 51)
store.level <- add.lag.column(store.level, 52)
store.level <- add.lag.column(store.level, 53)
store.level$Lag_1_53_Diff <- store.level$Lag_1_Sales - store.level$Lag_53_Sales

```

Some plots of other predictor variables we'll include in our models:


```
ggplot(data=store.level, aes(x=IsHoliday, y=Norm_Sales)) + geom_boxplot()
```



```

store.level <- add.normalized.col(store.level, 'Temperature')
ggplot(data=store.level, aes(x=Norm_Temperature, y=Norm_Sales)) + geom_point(alpha=1/5)

```



## Setup

We limit the testing/training sets to data after February 11, 2011, because we cannot calculate all of our lag features for weeks before then. We then divide all the remain data points randomly into an 80% training set and a 20% test set.

```
test.train <- subset(store.level, Date >= '2011-02-11')
set.seed(123)
row.count=nrow(test.train)
train.ind <- sample(1:row.count, size=row.count*0.8)
train.set <- test.train[train.ind, ]
test.set <- test.train[-train.ind, ]
paste(nrow(train.set), 'training records;', nrow(test.set), 'testing records')
```

```
## [1] "3240 training records; 810 testing records"
```

To allow us to make good comparisons between models, we use a consistent set of predictors for all the models we fit in the following section. We include:

- Normalized external economic factors (unemployment rate, CPI, and fuel price)
- A holiday indicator variable with a Store\_Type interaction (because we saw that store type C was less impacted by the holiday season than were the other store types)
- Normalized temperature
- Month-of-year and lag features

```
sales.formula <- Norm_Sales ~ Norm_Unemployment + Norm_CPI + Norm_Fuel_Price +
  IsHoliday*Store_Type + Norm_Temperature + Lag_1_Sales + Lag_2_Sales + Lag_3_Sales +
  Lag_4_Sales + Lag_51_Sales + Lag_52_Sales + Lag_1_53_Diff + Month
```

We define a convenience function for evaluating the model on the training and test sets:

```
get.mses <- function(model){
  test.mse <- mean((predict(model, test.set) - test.set$Norm_Sales)^2)
  train.mse <- mean((predict(model, train.set) - train.set$Norm_Sales)^2)
  paste('Train MSE:', train.mse, '; Test MSE:', test.mse)
}
```

## Model comparison

To start, we try a linear regression model. With no variable selection or shrinkage, it wouldn't be surprising to see the model overfit the data.

```
linear.model <- lm(sales.formula, data=train.set)
summary(linear.model)
```

```
##
## Call:
## lm(formula = sales.formula, data = train.set)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-3.09847	-0.23758	-0.01113	0.22817	2.30696

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-0.496979	0.055113	-9.018	< 2e-16 ***
Norm_Unemployment	-0.027831	0.011355	-2.451	0.014296 *
Norm_CPI	0.039837	0.016262	2.450	0.014352 *
Norm_Fuel_Price	0.002536	0.024391	0.104	0.917206
IsHolidayTRUE	0.308418	0.047848	6.446	1.32e-10 ***
Store_TypeB	-0.006876	0.017187	-0.400	0.689153
Store_TypeC	0.081380	0.024458	3.327	0.000887 ***
Norm_Temperature	0.031247	0.024125	1.295	0.195337
Lag_1_Sales	0.066701	0.012006	5.556	2.99e-08 ***
Lag_2_Sales	0.042111	0.010439	4.034	5.61e-05 ***
Lag_3_Sales	0.051329	0.011213	4.578	4.88e-06 ***
Lag_4_Sales	0.109297	0.011346	9.633	< 2e-16 ***
Lag_51_Sales	-0.097507	0.010285	-9.481	< 2e-16 ***
Lag_52_Sales	0.718462	0.010520	68.296	< 2e-16 ***
Lag_1_53_Diff	0.395572	0.014301	27.660	< 2e-16 ***
Month2	0.633779	0.057863	10.953	< 2e-16 ***
Month3	0.430245	0.054500	7.894	3.97e-15 ***
Month4	0.458819	0.063369	7.240	5.57e-13 ***
Month5	0.466137	0.068735	6.782	1.41e-11 ***
Month6	0.451954	0.074134	6.096	1.21e-09 ***
Month7	0.376220	0.077686	4.843	1.34e-06 ***
Month8	0.455013	0.078680	5.783	8.04e-09 ***
Month9	0.277200	0.071454	3.879	0.000107 ***
Month10	0.522266	0.062635	8.338	< 2e-16 ***
Month11	0.805508	0.062933	12.799	< 2e-16 ***
Month12	0.741737	0.065486	11.327	< 2e-16 ***
IsHolidayTRUE:Store_TypeB	0.066863	0.066245	1.009	0.312891
IsHolidayTRUE:Store_TypeC	-0.443131	0.093618	-4.733	2.30e-06 ***

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4347 on 3212 degrees of freedom
## Multiple R-squared:  0.774, Adjusted R-squared:  0.7721
## F-statistic: 407.4 on 27 and 3212 DF, p-value: < 2.2e-16
```

In the summary, we see that:

- Norm\_CPI appears very significant while the other economic indicators do not. However, we suspect there is some collinearity between these three predictors, so we do not want to read too much into their relative significances. It seems that collinearity might also be leading to the negative Norm\_Fuel\_Price coefficient; based on our earlier analysis, we expected this to be negative.
- The IsHolidays indicator variable has a very significant positive coefficient (i.e. holidays lead to more sales), and the store type C interaction has a significant negative coefficient (which aligns with our earlier observation that store type C did not exhibit the same seasonal behavior as the other store types in our timeseries plots).
- Of all the lag variables, the sales from 52 weeks ago seem most important for prediction. The Lag\_1\_53\_Diff variable, which we are using as a proxy for long-term trend, also stands out as being helpful for prediction.

How does the MSE on the training set compare to the MSE on the test set?

```
get.msres(linear.model)
```

```
## [1] "Train MSE: 0.187299169969781 ; Test MSE: 0.191814323096556"
```

The test error is only slightly higher than the training error. This seems to indicate that we do not have a notable overfittig problem, even though we have not used any subset selection or shrinkage. Below, we apply subset selection using forward stepwise regression to see if/how the results change.

```
# TODO: apply step
```


TODO: KNN, tree

TODO: comparison

## Appendix

Here are some interesting observations that didn't directly fit into the analysis presented in the body of the report.

```
sales.by.size <- get.aggr.sales(store.level, c('Store', 'Store_Type', 'Store_Size'))
ggplot(data=sales.by.size, aes(x=Store_Size, y=Sales_Millions, color=Store_Type)) +
  geom_point()
```



There is roughly a linear relationship between a store's sales and its size. We also notice in the plot above that store type is closely related to both store sales and store size.